



# Chapter 14: Transactions

**Database System Concepts, 7<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Transaction Concept

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items
- E.g., transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- Two main issues to deal with:
  - Failures of various kinds, such as hardware failures and system crashes
  - Concurrent execution of multiple transactions



# Example of Fund Transfer

- Transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- **Atomicity requirement**
  - If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
    - Failure could be due to software or hardware
  - The system should ensure that updates of a partially executed transaction are not reflected in the database
- **Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures



# Example of Fund Transfer

- **Consistency requirement** in above example:
  - The sum of A and B is unchanged by the execution of the transaction
- In general, consistency requirements include
  - Explicitly specified integrity constraints such as primary keys and foreign keys
  - Implicit integrity constraints
    - e.g., sum of balances of all accounts, minus sum of loan amounts must equal value of cash-in-hand
  - A transaction must see a consistent database
  - During transaction execution the database may be temporarily inconsistent
  - When the transaction completes successfully the database must be consistent
    - Erroneous transaction logic can lead to inconsistency



# Example of Fund Transfer

- **Isolation requirement** — if between steps 3 and 6, another transaction T2 is allowed to access the partially updated database, it will see an inconsistent database (the sum  $A + B$  will be less than it should be)

**T1**

1. **read**( $A$ )
2.  $A := A - 50$
3. **write**( $A$ )
4. **read**( $B$ )
5.  $B := B + 50$
6. **write**( $B$ )

**T2**

read( $A$ ), read( $B$ ), print( $A+B$ )

- Isolation can be ensured trivially by running transactions **serially**
  - That is, one after the other
- However, executing multiple transactions concurrently has significant benefits



# ACID Properties

A **transaction** is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:

- **Atomicity.** Either all operations of the transaction are properly reflected in the database, or none are (all or nothing)
- **Consistency.** Execution of a transaction in isolation preserves the consistency of the database
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
  - That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures




# Transaction State

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed
- **Failed** -- after the discovery that normal execution can no longer proceed
- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
  - Restart the transaction
    - Can be done only if no internal logical error
  - Kill the transaction
- **Committed** – after successful completion



# Transaction State







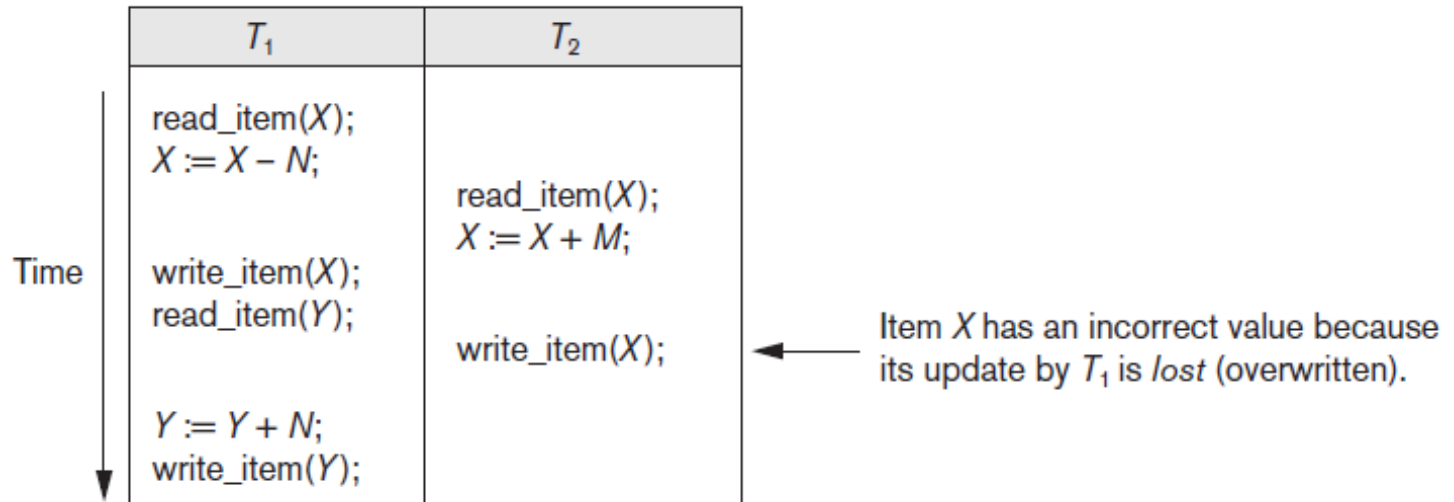
# Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system  
Advantages are:
  - **Increased processor and disk utilization**, leading to better transaction *throughput*
    - E.g., one transaction can be using the CPU while another is reading from or writing to the disk
  - **Reduced average response time** for transactions: short transactions need not wait behind long ones
- **Concurrency control schemes** – mechanisms to achieve isolation
  - That is, to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database



# The Lost Update Problem

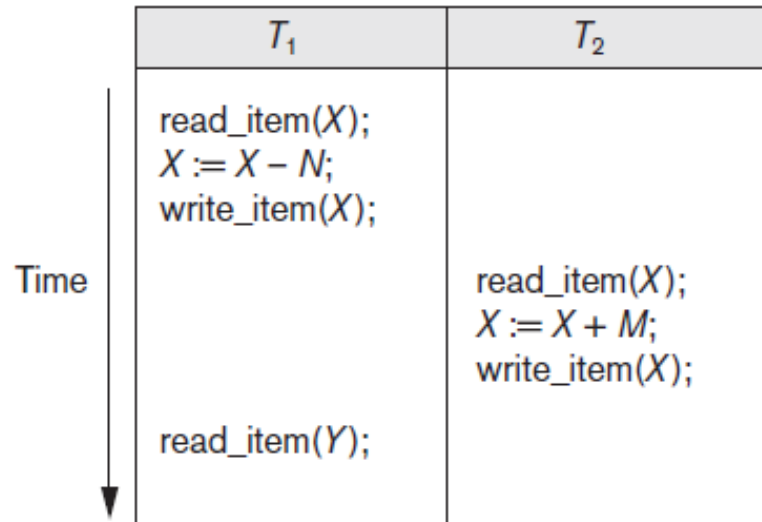
(a)





# The Temporary Update Problem

(b)



Transaction  $T_1$  fails and must change the value of  $X$  back to its old value; meanwhile  $T_2$  has read the *temporary* incorrect value of  $X$ .



# The Incorrect Summary Problem

(c)

$T_1$	$T_3$
<pre>read_item(X); X := X - N; write_item(X);  read_item(Y); Y := Y + N; write_item(Y);</pre>	<pre>sum := 0; read_item(A); sum := sum + A; . . .  read_item(X); sum := sum + X; read_item(Y); sum := sum + Y;</pre>

←  $T_3$  reads  $X$  after  $N$  is subtracted and reads  $Y$  before  $N$  is added; a wrong summary is the result (off by  $N$ ).



# The Unrepeatable Read Problem

- Transaction T reads the same item twice
- Value is changed by another transaction T' between the two reads
- T receives different values for the two reads of the same item



# Schedules

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - A schedule for a set of transactions must consist of all instructions of those transactions
  - Must preserve the order in which the instructions appear in each individual transaction
- A transaction that successfully completes its execution will have a commit instructions as the last statement
  - By default, a transaction is assumed to execute commit instruction as its last step



# Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance of  $A$  from  $A$  to  $B$
- A **serial** schedule in which  $T_1$  is followed by  $T_2$  :

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit



## Schedule 2

- A serial schedule where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	





# Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1 (i.e.,  $T_1$  followed by  $T_2$ )

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ )	
	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ )
read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	
	read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

- In Schedules 1, 2 and 3, the sum  $A+B$  is preserved



- | $T_1$   | $T_2$   |
|---|---|
| read ( $A$ )<br>$A := A - 50$   | read ( $A$ )<br>$temp := A * 0.1$<br>$A := A - temp$<br>write ( $A$ )<br>read ( $B$ ) |
| write ( $A$ )<br>read ( $B$ )<br>$B := B + 50$<br>write ( $B$ )<br>commit | $B := B + temp$<br>write ( $B$ )<br>commit  |



# Serializability

- **Basic Assumption** – Each transaction is assumed correct if executed on its own
- Serial execution of a set of transactions is assumed correct
- **Criterion for correctness: every serial schedule is considered correct**
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notion of **conflict serializability**



# Simplified view of transactions

- We ignore operations other than **read** and **write** instructions
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes
- Our simplified schedules consist of only **read** and **write** instructions



# Conflicting Instructions

- Instructions  $I$  and  $J$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists some item  $Q$  accessed by both  $I$  and  $J$ , and at least one of these is a write instruction
  1.  $I = \text{read}(Q)$ ,  $J = \text{read}(Q)$ .  $I$  and  $J$  don't conflict (the order of  $I$  &  $J$  does not matter)
  2.  $I = \text{read}(Q)$ ,  $J = \text{write}(Q)$ . They conflict (the order of  $I$  &  $J$  matters: write before read and read before write gives different read results)
  3.  $I = \text{write}(Q)$ ,  $J = \text{read}(Q)$ . They conflict (the order of  $I$  &  $J$  matters)
  4.  $I = \text{write}(Q)$ ,  $J = \text{write}(Q)$ . They conflict (the order of  $I$  &  $J$  matters since it would affect the result of the next **read**( $Q$ ) instruction)
- Intuitively, a conflict between  $I$  and  $J$  forces a (logical) temporal order between them
- If  $I$  and  $J$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule



# Conflict Serializability

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**
- Or equivalently, two schedules are conflict equivalent if the relative order of any two conflicting instructions is the same in both schedules
- We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule



# Conflict Serializability

- Schedule 1 can be transformed into Schedule 2, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions. Therefore Schedule 1 is conflict serializable

$T_1$	$T_2$
read (A) write (A)	read (A) write (A)
read (B) write (B)	
	read (B) write (B)

Schedule 1

$T_1$	$T_2$
read (A) write (A) read (B) write (B)	read (A) write (A) read (B) write (B)

Schedule 2



# Conflict Serializability

- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read ( $Q$ )	write ( $Q$ )
write ( $Q$ )	


- $T_4$ 's update is lost
- We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$
- This is not conflict serializable since it is not equivalent to either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$





# Testing for Serializability




- Consider some schedule  $S$  of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph** — a directed graph where the vertices are the transactions of a schedule  $S$
- We draw an arc from  $T_i$  to  $T_j$  if one of three conditions holds:
  - $T_i$  executes **write**( $Q$ ) before  $T_j$  executes **read**( $Q$ )
  - $T_i$  executes **read**( $Q$ ) before  $T_j$  executes **write**( $Q$ )
  - $T_i$  executes **write**( $Q$ ) before  $T_j$  executes **write**( $Q$ )
- If an edge  $T_i \rightarrow T_j$  exists in the precedence graph, then any serial schedule  $S'$  equivalent to  $S$ ,  $T_i$  must appear before  $T_j$
- Example of a precedence graph





# Test for Conflict Serializability

- A schedule is conflict serializable if and only if its precedence graph is acyclic
- Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph
- If precedence graph is acyclic, the serializability order can be obtained by a topological sorting of the graph
  - This is a linear order consistent with the partial order of the graph
  - A serializability order for Schedule (a) would be
$$T_i \rightarrow T_j \rightarrow T_k \rightarrow T_m$$





# Recoverable Schedules

Need to address the effect of transaction failures on concurrently running transactions.

- **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .
- The following schedule is not recoverable

$T_8$	$T_9$
read (A) write (A)	
	read (A) commit
read (B)	

- If  $T_8$  should abort,  $T_9$  would have read an inconsistent database state but  $T_9$  has already committed. Hence, database must ensure that schedules are recoverable



# Cascading Rollbacks

- **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks. Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	
abort		read (A)

If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back. The read (A) in  $T_{11}$  and  $T_{12}$  is called **dirty read**

- Can lead to the undoing of a significant amount of work



# Cascadeless Schedules

- **Cascadeless schedules** — cascading rollbacks cannot occur;
  - For each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$
- Every Cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless



# Weak Levels of Consistency

- Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
  - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
  - Such transactions need not be serializable with respect to other transactions
- Trade accuracy for performance



# Transaction Definition in SQL

- In SQL, a transaction begins implicitly
- A transaction in SQL ends by:
  - **Commit work** commits current transaction and begins a new one
  - **Rollback work** causes current transaction to abort



# Transaction Support in SQL

## ■ Isolation levels

- **Dirty read** (reading of the update of an uncommitted transaction)
- **Nonrepeatable read** (another transaction updates a data item between two reads so that the transaction sees two different values)
- **Phantoms** (if another transaction inserts a new record  $r$  during the execution of the transaction,  $r$  was not there at the beginning of the transaction but was there at the end of the transaction;  $r$  is called a **phantom record**)

Isolation Level	Type of Violation		
	Dirty Read	Nonrepeatable Read	Phantom
READ UNCOMMITTED	Yes	Yes	Yes
READ COMMITTED	No	Yes	Yes
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No